Project 1 document
Part 1:

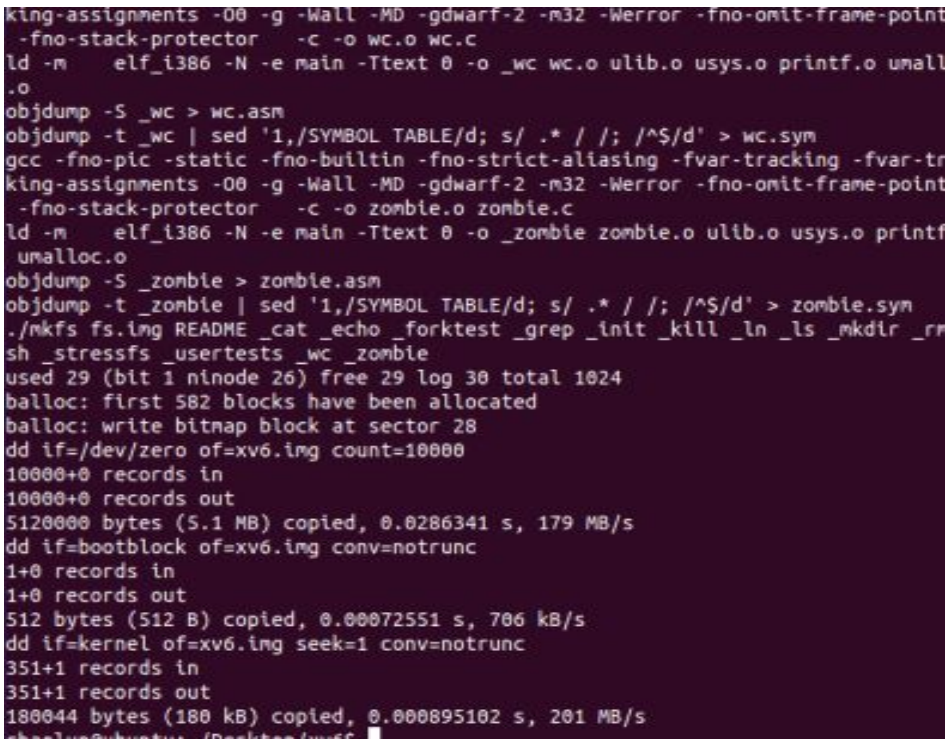

Successfully compiled the xv6 by typing make in linux terminal

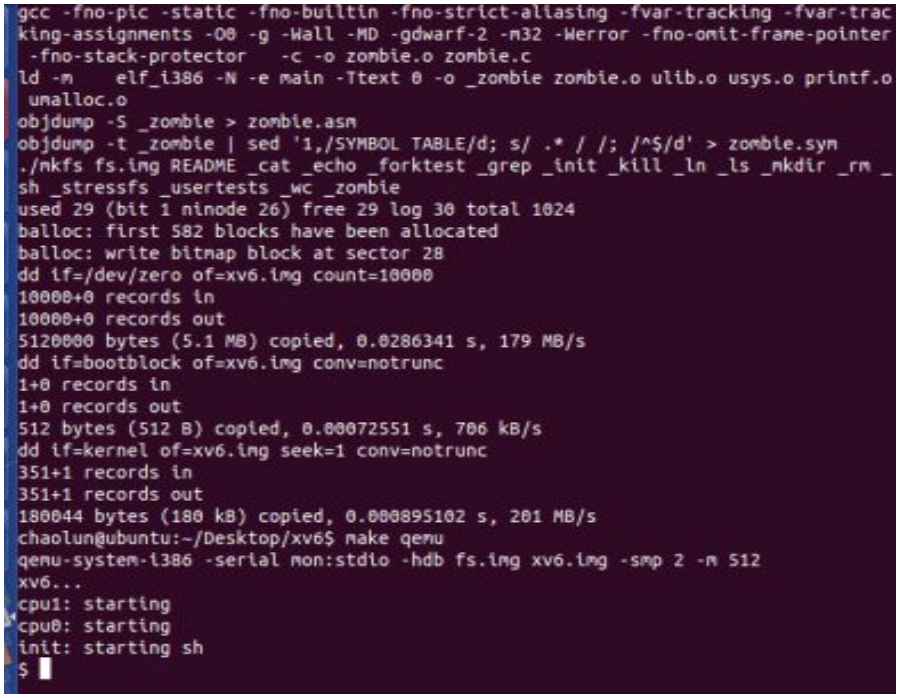

Successfully run the xv6 on QEMU

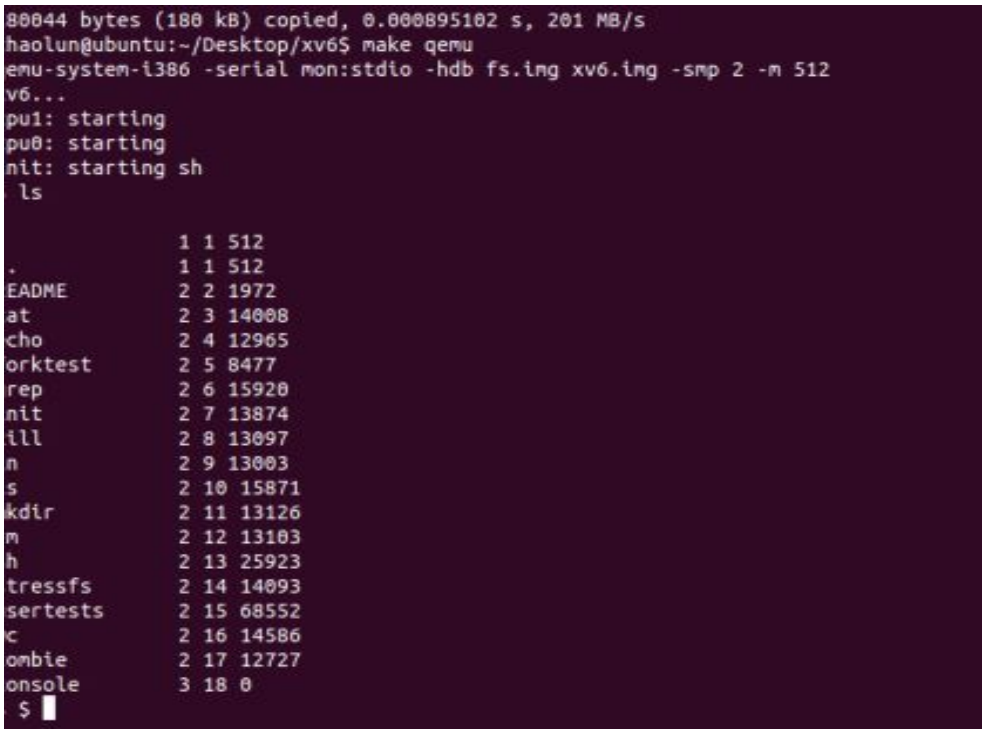

Using ls command to see if the xv6 is working well on QEMU

Part 2
Analyzing read(10, buf, n)

In user space:

1) read function is declared in user.h, line 10 :
   int read(int, void*, int);
2) read function is defined in usys.S ,line 4 - line 9 :
   SYSCALL(read), which was defined as:
   .globl read; //declares read as a global symbol
   read: //which is the entry point of read
   movl $SYS_read, %eax; //store system call number in eax register
   int $T SYSCALL; //trigger a software interrupt and enter the kernel
   ret //return result to the caller of read
3) $SYS_read was defined in syscall.h, line 6
   #define SYS_read 5 //system call number of read is set as 5
4) $T_SYSCALL was defined in traps.h, line 27
   #define $T_SYSCALL 64 // 64 is also the No. for the interrupt handler vector

Enter the Kernal:
int $T_SYSCALL triggers a software interrupt:

1) CPU saves the current state, and calls the interrupt handler vector64.

2) Vector64 is in vectors.S, line317-321:

```
.globl vector64
vector64:
      pushl $0
      pushl $64
      jmp alltraps                        //jump to alltraps function
```

3) vector64 jump to alltraps function in vectors.S, line 321.

4) alltraps creates the trapframe and calls trap(struct trapframe *tf) function, which is in file trap.c line 37-47

5) struct trapframe (tf) saves the user-space registers, tf->eax contains the system call number which is SYS_read.

```
proc->tf = tf;          //line42 in file trap.c, load the value from registers
```

Syscall dispatch:

1)trap (tf), which is in trap.c, calls syscall (void) : //happens in line43 in file trap.c

```
void
trap(struct trapframe *tf)
{
  if(tf->trapno == T_SYSCALL)          //will return true in this case
  {
    if(proc->killed)                   //check if the process have been killed.
      exit();                          //if true exit
    proc->tf = tf;                   //load value form tf passing into trap function
    syscall();                         //call syscall function
    if(proc->killed)                   //check if the process have been killed.
      exit();                          //if true exit
    return;
  }
}
```

2)syscall, which is in syscall.c line126-139, reads the syscall number in eax, and calls sys_read:

```
void
syscall(void)
{
  int num;
```

```
    num = proc->tf->eax;                                    //read syscall number in eax
    if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
        proc->tf->eax = syscalls[num]();                    //return is saved in tf->eax

    } else {
        cprintf("%d %s: unknown sys call %d\n",
                    proc->pid, proc->name, num);
        proc->tf->eax = -1;                                 //unknown sys_call number return -1
    }
}
```

3)sys_read is defined to be syscalls[SYS_read], in file syscall.c line 107:

```
static int (*syscalls[])(void) = {
          .....
     [SYS_read]       sys_read,            // //syscalls[read] is defined to be sys_read
         ......
}
```

4)sys_read, which is in sysfile.c line66-76, reads the parameter from user stack with argfd, argint, and argptr. Since fd=10 is invalid, argfd will return -1, so the return value of sys_read is -1.

```
int
sys_read(void)
{
    struct file *f;
    int n;
    char *p;

    if(argfd(0, 0, &f) < 0 || argint(2, &n) < 0 || argptr(1, &p, n) < 0)
        return -1;                          //In this case,argfd(0, 0, &f) < 0 ==True, return -1
    return fileread(f, p, n);           //if sys_read() is valid, read form file
}
```

5)return value is saved in tf->eax(in syscall.c line132), control was then turned back to trap:

```
     proc->tf->eax = syscalls[num]();    //syscalls[read]==-1, saved to        proc->tf->eax
```

6)trap returns to alltraps, restores user registers and returns to user space with iret.